



CS 4650/7650: Natural Language Processing

Machine Translation I

Diyi Yang

Slides credit to Yulia Tsvetkov (CMU)

Machine Translation

A large, dense word cloud centered around the word "hello" in various languages. The words are in different colors and sizes, creating a visual representation of global communication. Some prominent words include "hello" in English, "bonjour" in French, "hallo" in German, "안녕하세요" in Korean, "こんにちは" in Japanese, and "привет" in Russian.



Tower of Babel

Text

Documents

DETECT LANGUAGE

ENGLISH

SPANISH

FRENCH

^



ENGLISH

SPANISH

ARABIC

▼

← Search languages

 Detect language

Danish

Hmong

Lithuanian

Romanian

Telugu

Afrikaans

Dutch

Hungarian

Luxembourgish

Russian

Thai

Albanian

English

Icelandic

Macedonian

Samoan

Turkish

Amharic

Esperanto

Igbo

Malagasy

Scots Gaelic

Turkmen

Arabic

Estonian

Indonesian

Malay

Serbian

Ukrainian

Armenian

Filipino

Irish

Malayalam

Sesotho

Urdu

Azerbaijani

Finnish

Italian

Maltese

Shona

Uyghur

Basque

French

Japanese

Maori

Sindhi

Uzbek

Belarusian

Frissian

Javanese

Marathi

Sinhala

Vietnamese

Bengali

Galician

Kannada

Mongolian

Slovak

Welsh

Bosnian

Georgian

Kazakh

Myanmar (Burmese)

Slovenian

Xhosa

Bulgarian

German

Khmer

Nepali

Somali

Yiddish

Catalan

Greek

Kinyarwanda

Norwegian

Spanish

Yoruba

Cebuano

Gujarati

Korean

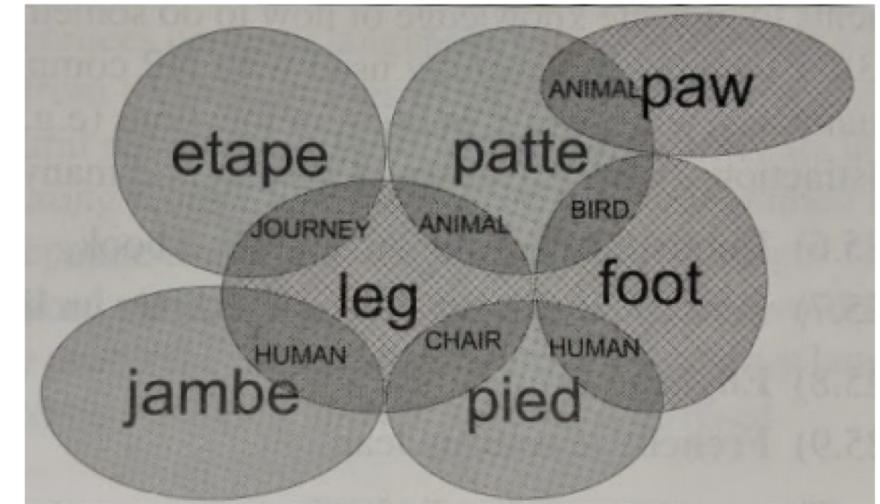
Odia (Oriya)

Sundanese

Zulu

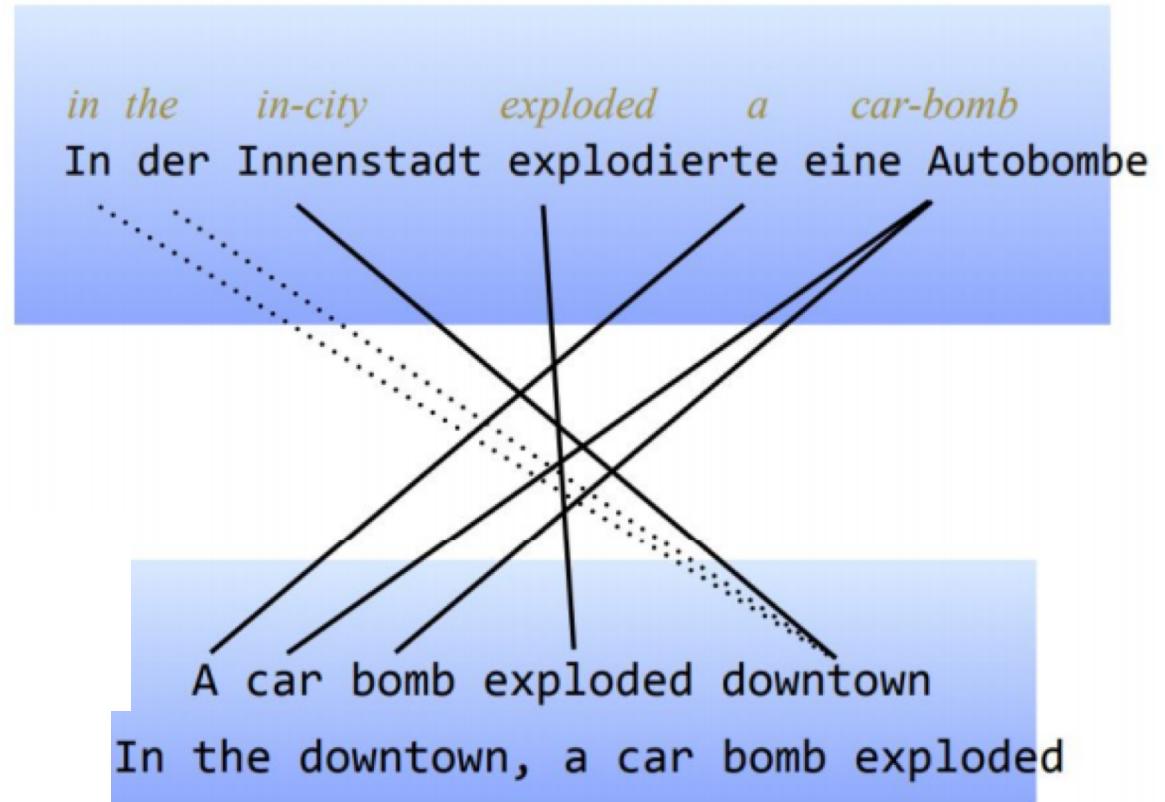
Dictionaries

- English: leg, foot, paw
- French: jambe, pied, patte, etape



Challenges

- Ambiguities
 - Words
 - Morphology
 - Syntax
 - Semantics
 - Pragmatics
- Gaps in data
 - Availability of corpus
 - Commonsense knowledge
- Understanding of context, connotation, social norms, etc



Research Problems

- How can we formalize the process of learning to translate from examples?
- How can we formalize the process of finding translations for new inputs?
- If our model produces many outputs, how do we find the best one?
- If we have a gold standard translation, how can we tell if our output is good or bad?

Two Views Of MT

MT as Code Breaking

One naturally wonders if the problem of translation could conceivably be treated as a problem in cryptography. When I look at an article in Russian, I say: '*This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.*'



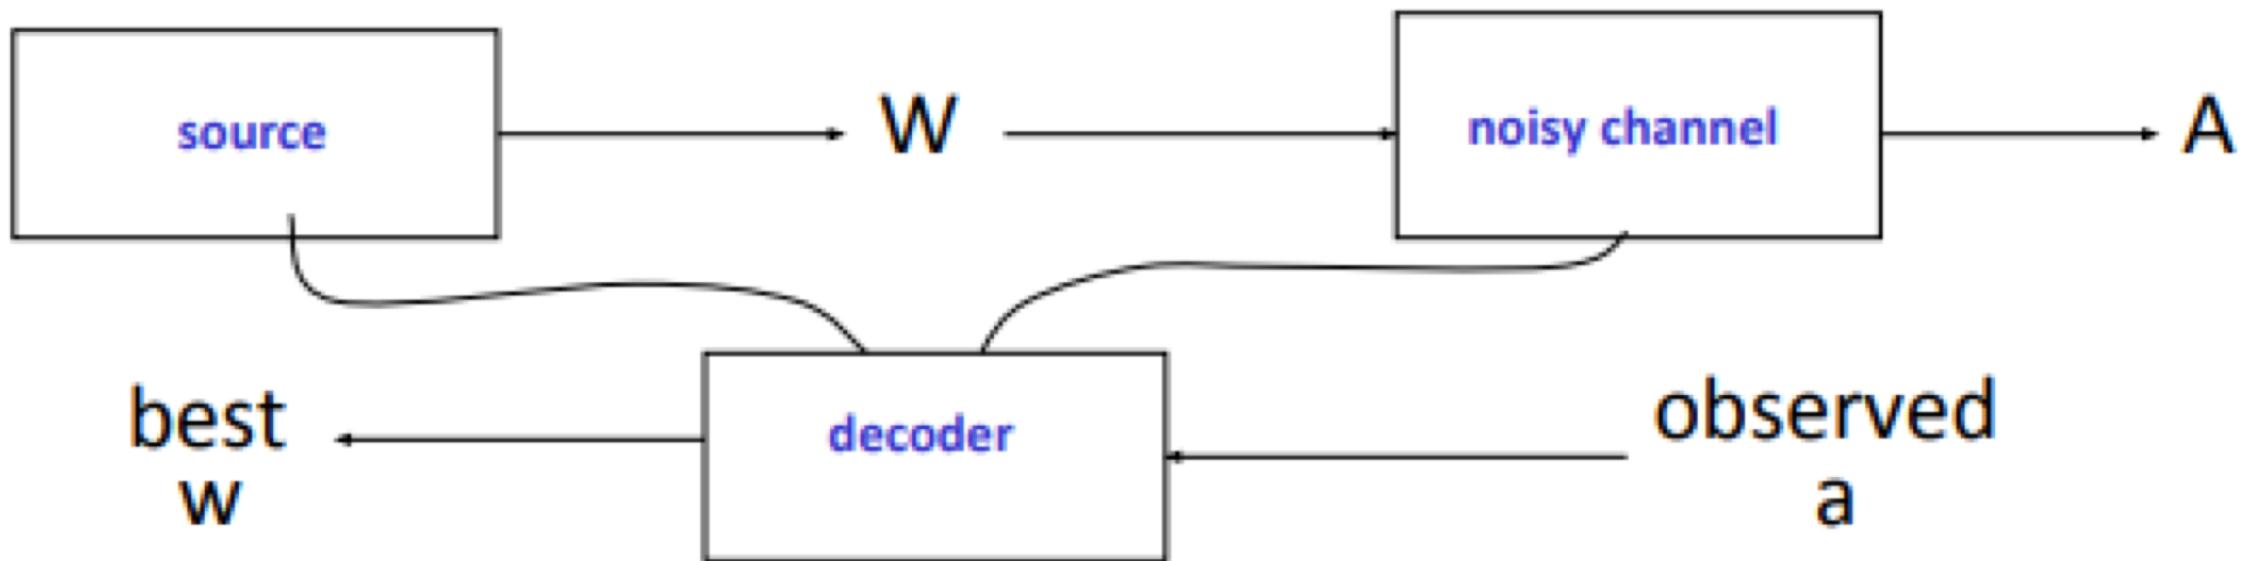
Warren Weaver to Norbert Wiener, March, 1947

The Noisy-Channel Model

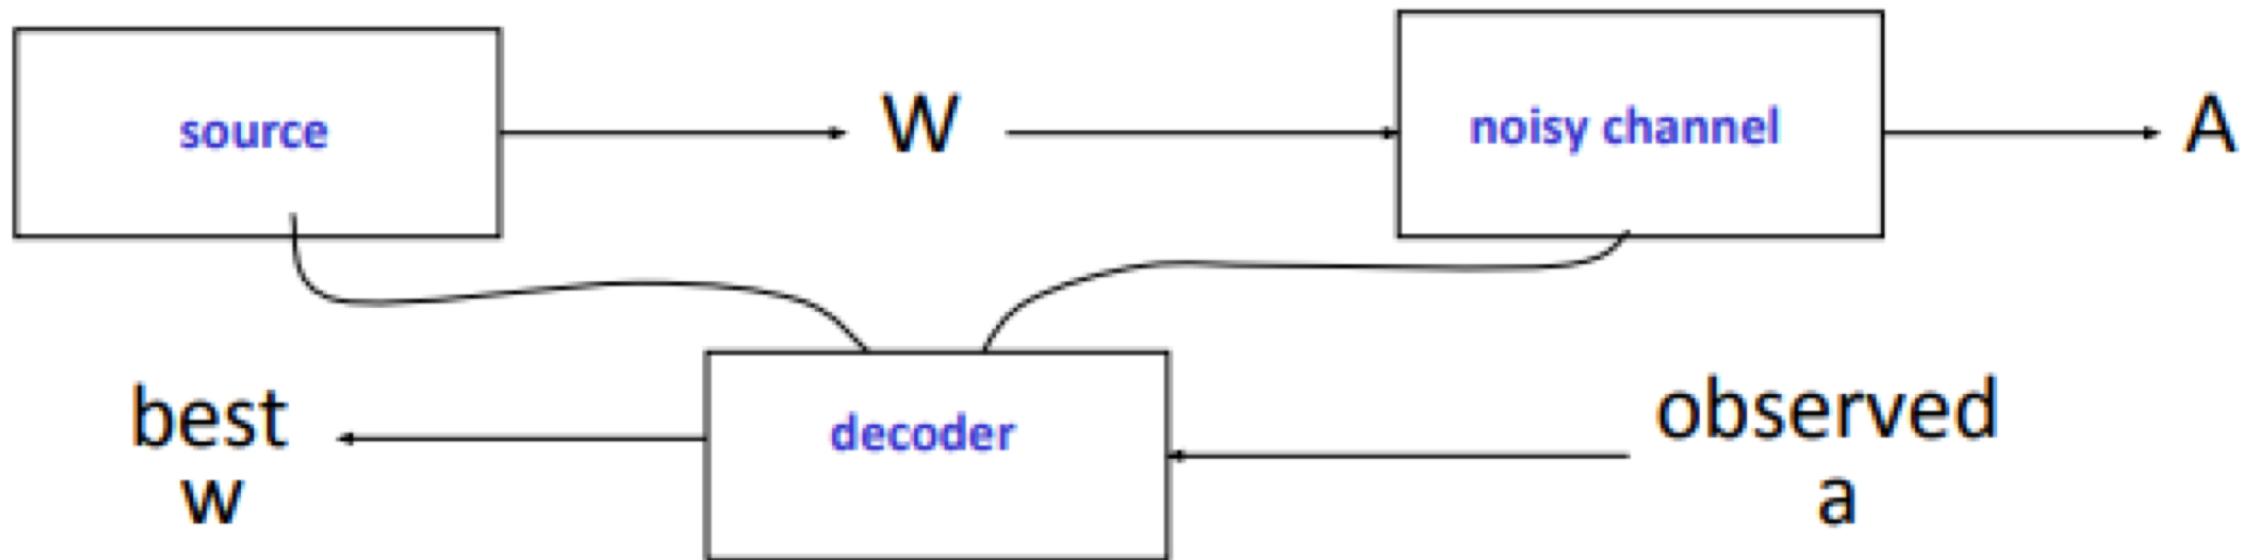


Claude Shannon. "A Mathematical Theory of Communication" 1948.

The Noisy-Channel Model



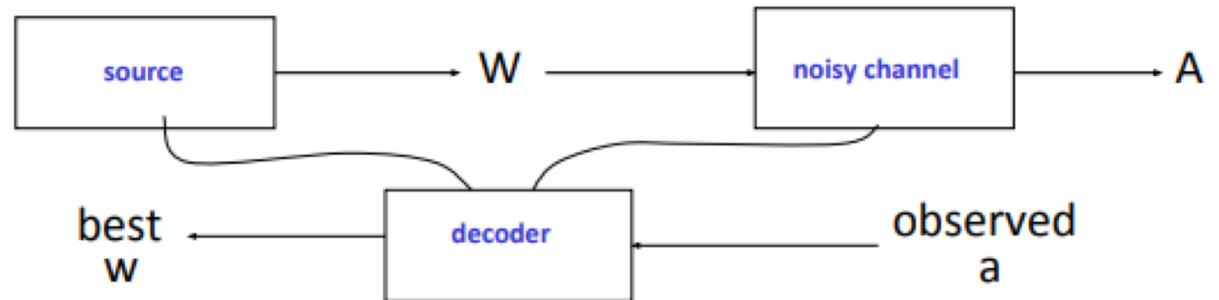
The Noisy-Channel Model



We want to predict a sentence given acoustics:

$$w^* = \arg \max_w P(w|a)$$

The Noisy-Channel Model



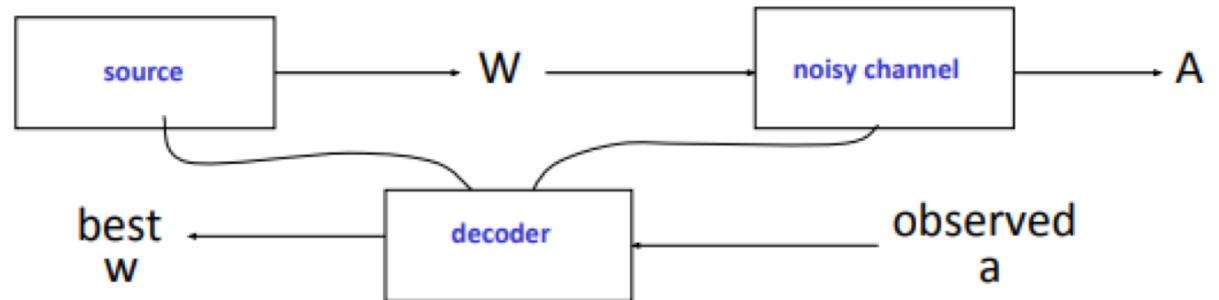
$$\begin{aligned}w^* &= \arg \max_w P(w|a) \\&= \arg \max_w \mathbf{P}(a|w) \mathbf{P}(w) / P(a)\end{aligned}$$

$$= \arg \max_w \mathbf{P}(a|w) \mathbf{P}(w)$$

Channel model

Source model

The Noisy-Channel Model



$$\begin{aligned}w^* &= \arg \max_w P(w|a) \\&= \arg \max_w \mathbf{P}(a|w) \mathbf{P}(w) / P(a) \\&= \arg \max_w \mathbf{P}(a|w) \mathbf{P}(w)\end{aligned}$$

Likelihood

Acoustic model (HMMs)
Translation model

Prior

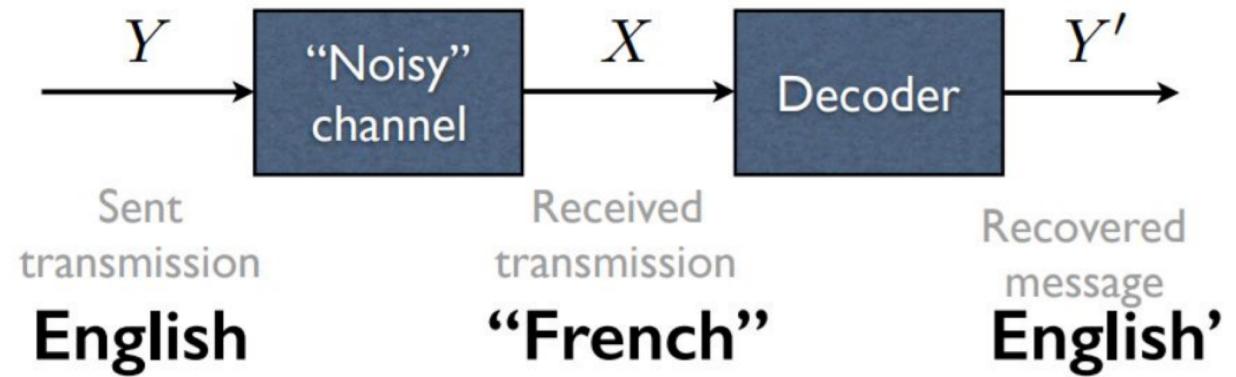
Language model: Distributions
over sequence of words

The Noisy-Channel Model

$$\hat{e} = \arg \max_e p_\varphi(e) \times p_\theta(f | e)$$

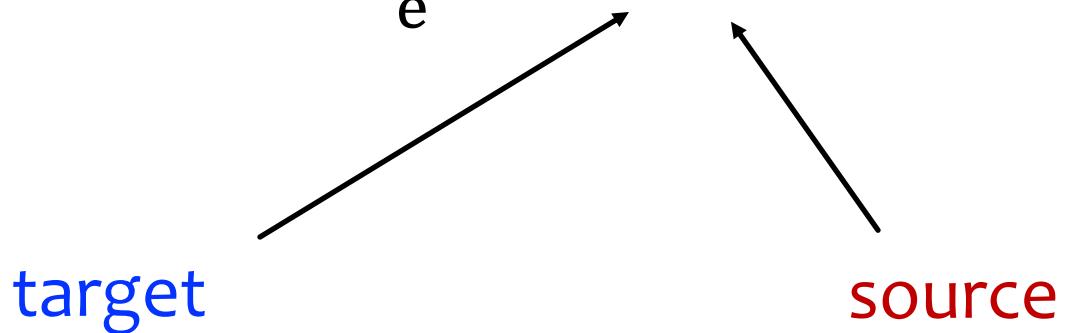
Language model

Translation model



MT as Direct Modeling

$$\hat{e} = \arg \max_e p_\lambda(e | f)$$



- One model does everything
- Trained to reproduce a corpus of translations

Two Views of MT

- **Code breaking** (aka the noisy channel, Bayes rule)
 - I know the **target language**
 - I have example **translations texts** (example enciphered data)
- **Direct modeling** (aka pattern matching)
 - I have **really good learning algorithms** and a bunch of **example inputs** (source language sentences) and **outputs** (target language translations)

Which is Better?

- **Noisy channel** - $p_\phi(e) \times p_\theta(f | e)$
 - Easy to use monolingual target language data
 - Search happens under a product of two models (individual models can be simple, product can be powerful)
- **Direct Model** - $p_\lambda(e | f)$
 - Directly model the process you care about
 - Model must be very powerful

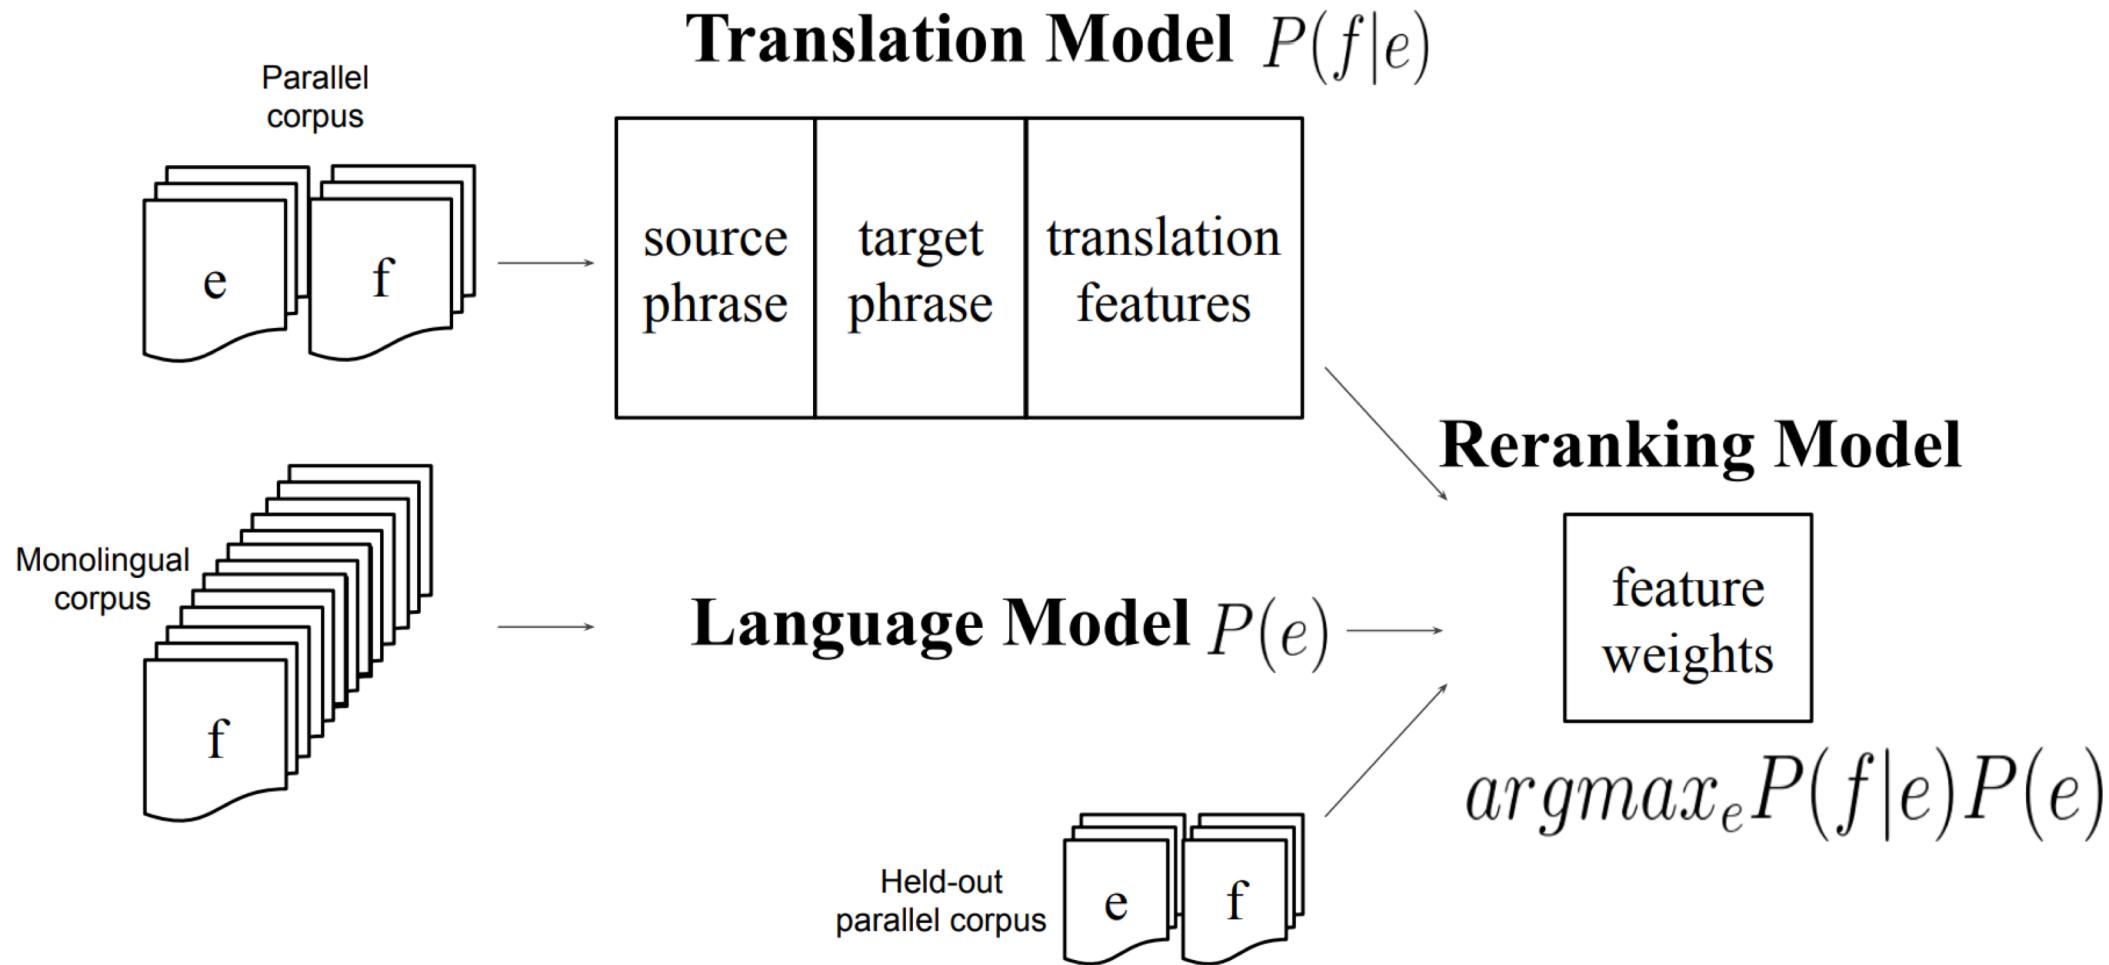
Where are we in 2020?

- Direct modeling is where most of the action is
 - Neural networks are very good at generalizing and conceptually very simple
 - Inference in “product of two models” is hard
- Noisy channel ideas are incredibly important and still play a big role in how we think about translation

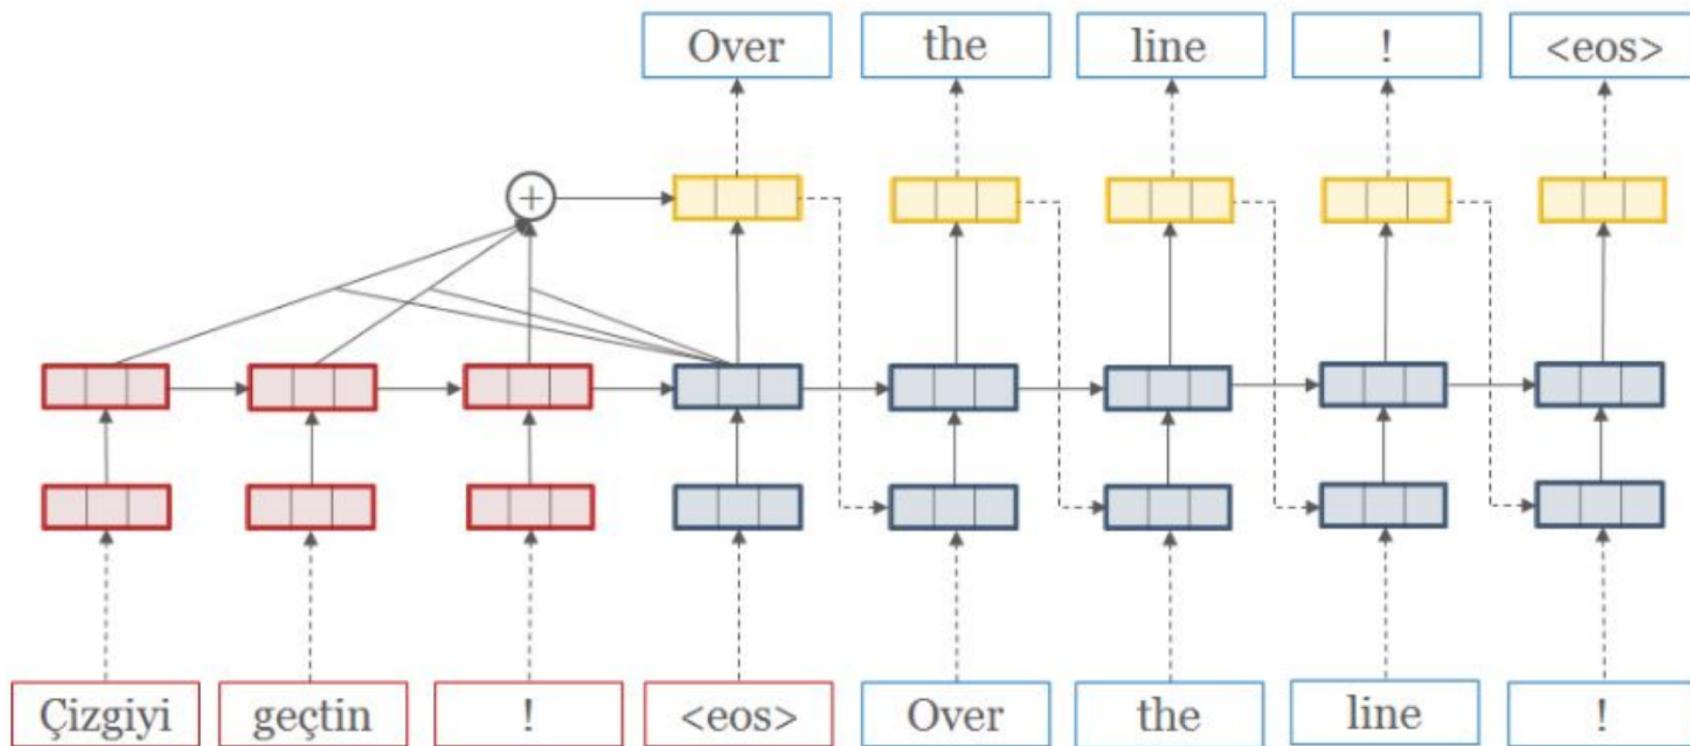
Two Views of MT

- Noisy channel $\hat{e} = \arg \max_e p_\varphi(e) \times p_\theta(f | e)$
- Direct $\hat{e} = \arg \max_e p_\lambda(e | f)$

Noisy Channel: Phrase-Based MT



Neural MT: Conditional Language Modeling

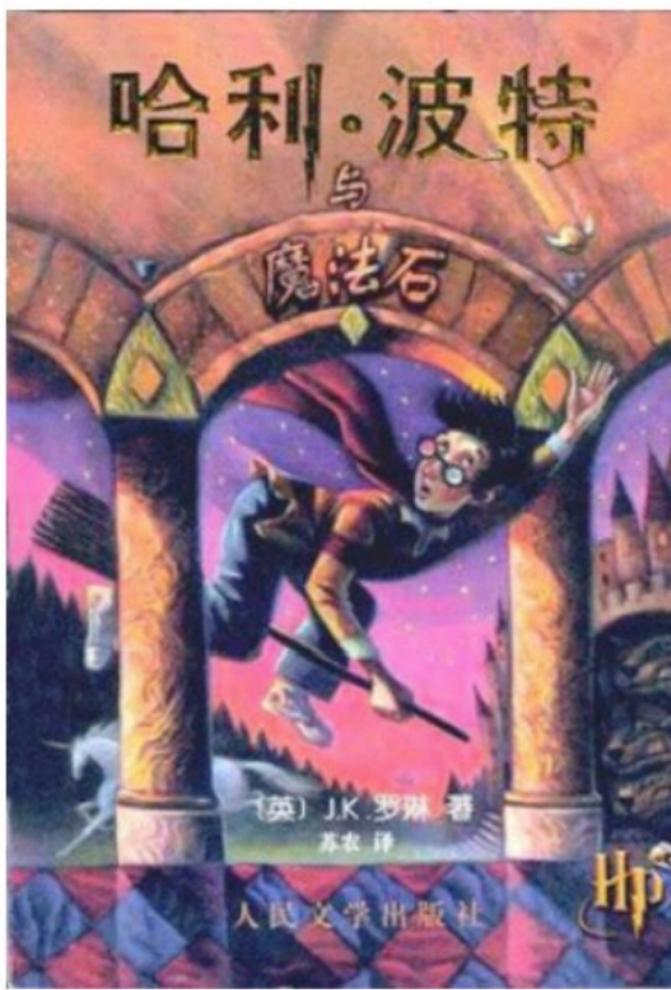
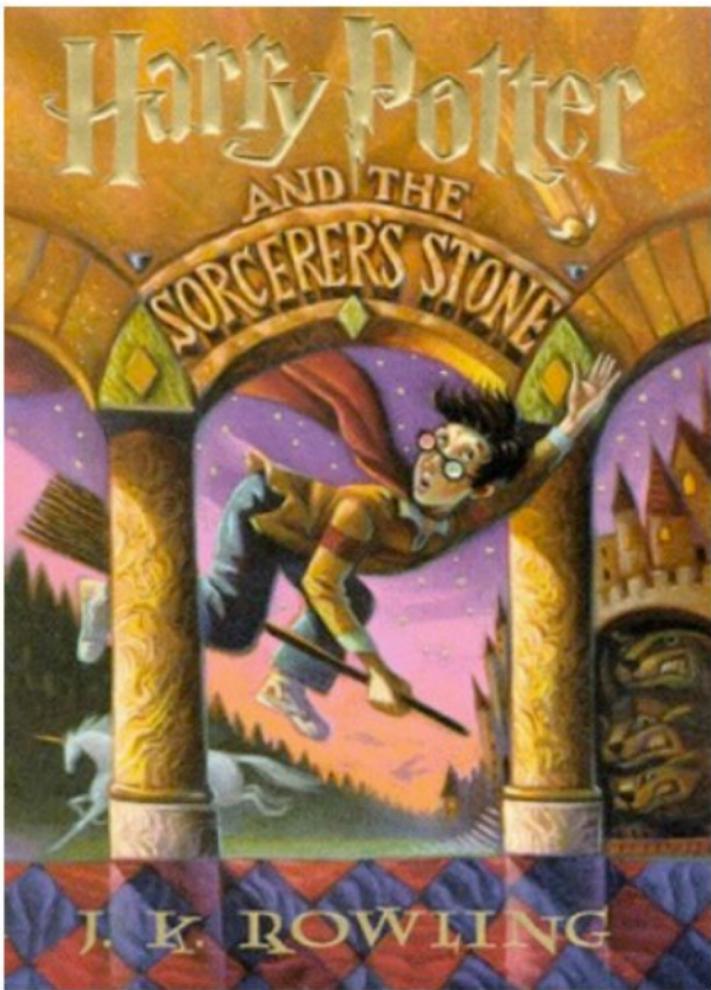


A Common Problem

- Noisy channel $\hat{e} = \arg \max_e p_\varphi(e) \times \textcolor{red}{p_\theta(f | e)}$
- Direct $\hat{e} = \arg \max_e \textcolor{red}{p_\lambda(e | f)}$
- Both models must assign probabilities to how a sentence in one language translates into a sentence in another language

Learning From Data

Parallel Corpora



Parallel Corpora

CLASSIC SOUPS

Sm. Lg.

清 燉 雞 湯	57.	House Chicken Soup (Chicken, Celery, Potato, Onion, Carrot)	1.50	2.75
雞 飯 湯	58.	Chicken Rice Soup	1.85	3.25
雞 麵 湯	59.	Chicken Noodle Soup	1.85	3.25
廣 東 雲 吞	60.	Cantonese Wonton Soup.....	1.50	2.75
蕃 茄 蛋 湯	61.	Tomato Clear Egg Drop Soup	1.65	2.95
雲 吞 湯	62.	Regular Wonton Soup	1.10	2.10
酸 辣 湯	63.	Hot & Sour Soup	1.10	2.10
蛋 花 湯	64.	Egg Drop Soup.....	1.10	2.10
雲 蛋 湯	65.	Egg Drop Wonton Mix.....	1.10	2.10
豆 腐 菜 湯	66.	Tofu Vegetable Soup	NA	3.50
雞 玉 米 湯	67.	Chicken Corn Cream Soup	NA	3.50
蟹 肉 玉 米 湯	68.	Crab Meat Corn Cream Soup.....	NA	3.50
海 鮮 湯	69.	Seafood Soup.....	NA	3.50

Parallel Corpora (mining parallel data from microblogs Ling et al., 2013)

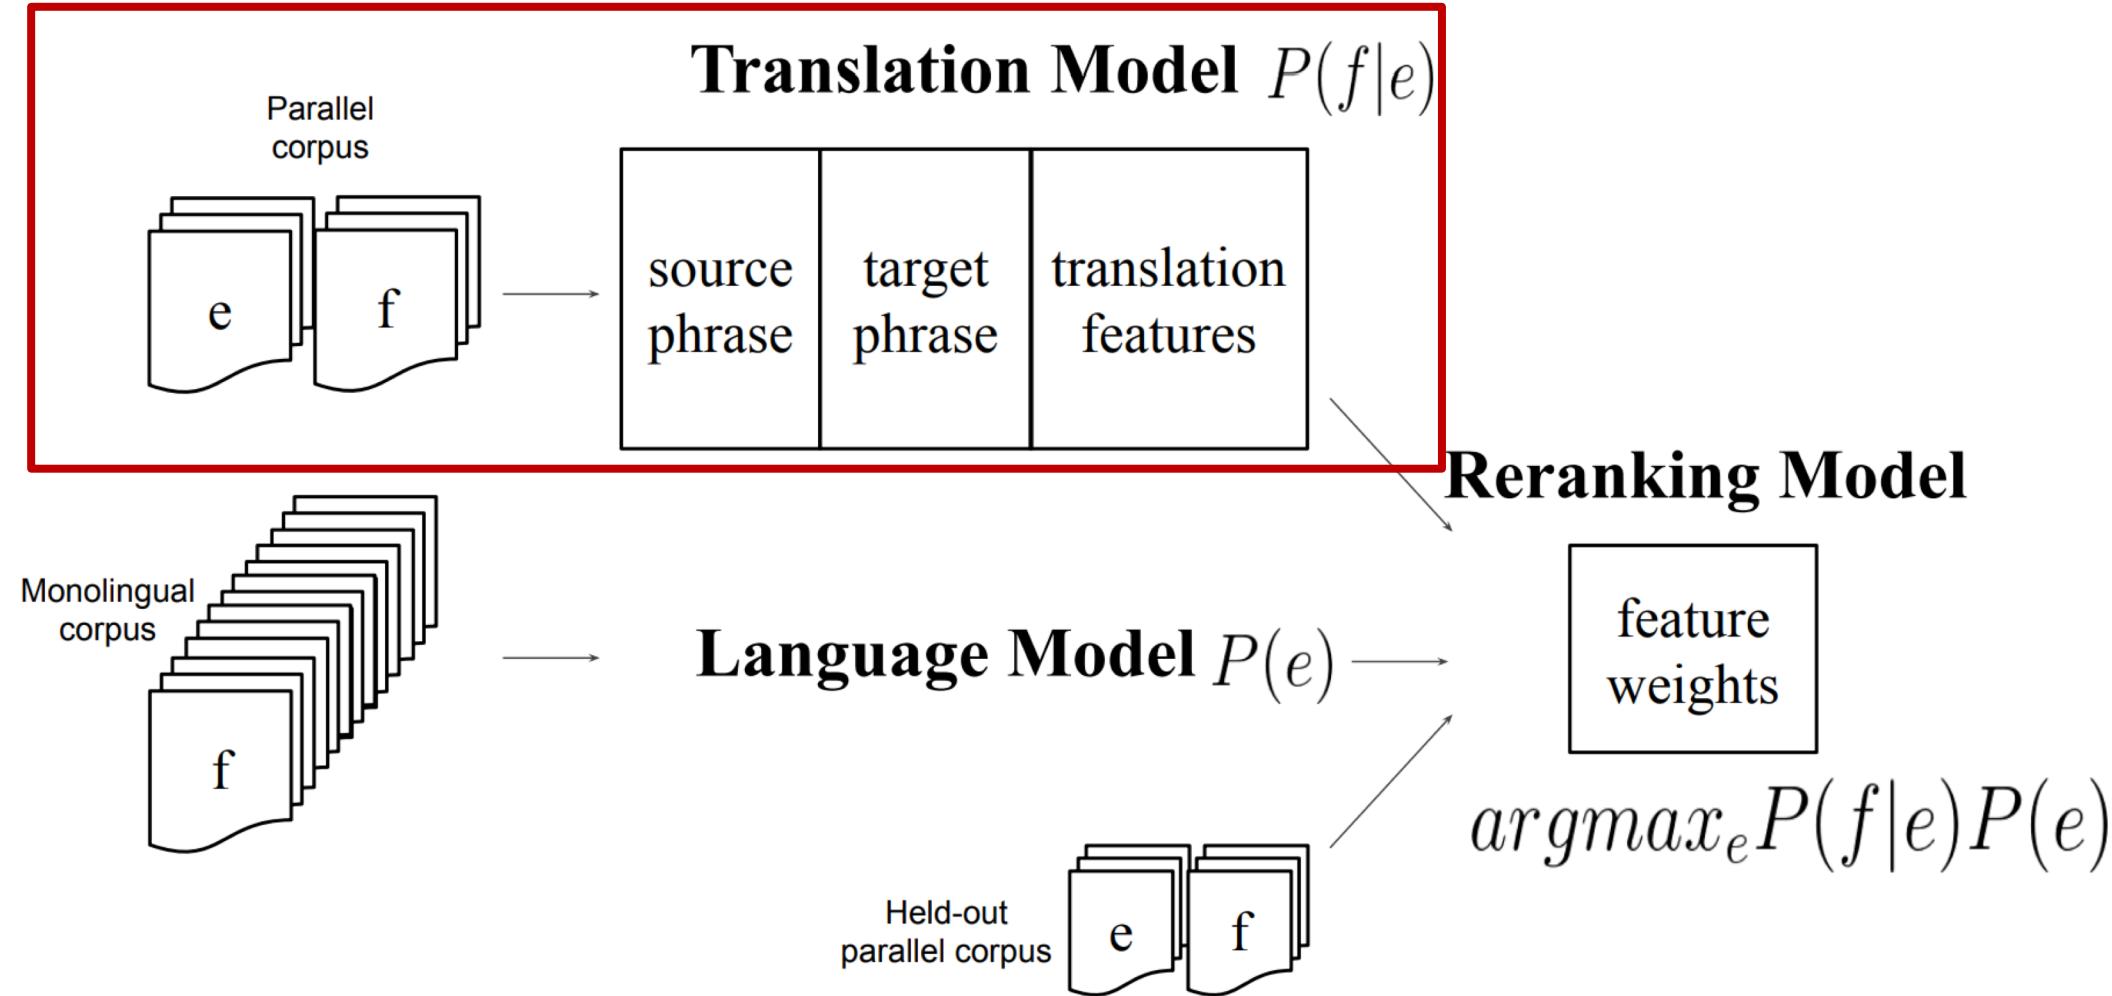
	ENGLISH	MANDARIN
1	i wanna live in a wes anderson world	我想要生活在Wes Anderson的世界里
2	Chicken soup, corn never truly digests. TMI .	鸡汤吧，玉米神马的从来没有真正消化过.恶心
3	To DanielVeuleman yea iknw imma work on that	对DanielVeuleman说，是的我知道，我正在向那方面努力
4	msg 4 Warren G his eday is today 1 yr older.	发信息给Warren G, 今天是他的生日，又老了一岁了。
5	Where the hell have you been all these years?	这些年你 TMD 到哪去了
	ENGLISH	ARABIC
6	It's gonna be a warm week!	الاسبوع اليابي حر
7	onni this gift only 4 u	أوني هذه الهدية فقط لك
8	sunset in aqaba :)	غروب الشمس في العقبة:)
9	RT @MARYAMALKHAWAJA: there is a call for widespread protests in #bahrain tmrw	هناك نداء لظاهرات في عدة مناطق غدا

Table 2: Examples of English-Mandarin and English-Arabic sentence pairs. The English-Mandarin sentences were extracted from Sina Weibo and the English-Arabic sentences were extracted from Twitter. Some messages have been shorted to fit into the table. Some interesting aspects of these sentence pairs are marked in bold.

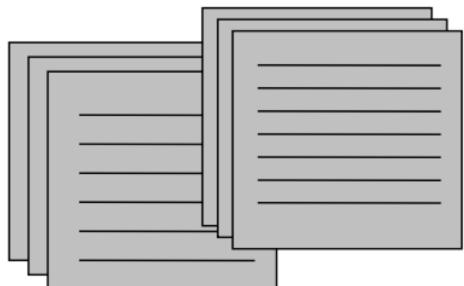
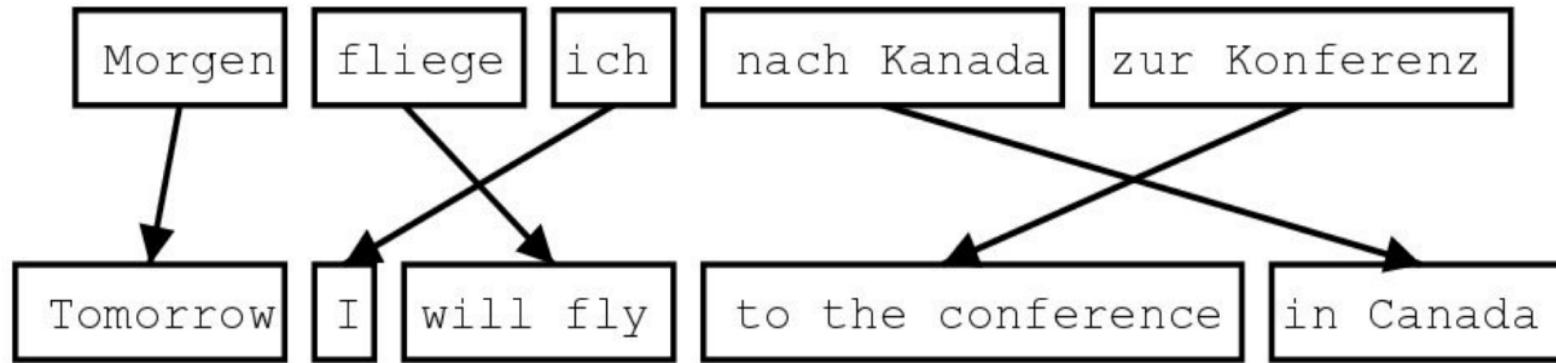
Discussions

- There is a lot more monolingual data in the world than translated data
- Easy to get about 1 trillion words of English by crawling the web
- With some work, you can get 1 billion translated words of English-French
 - What about Japanese-Turkish?

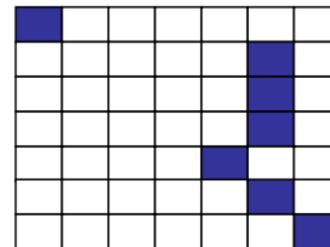
Phrase-Based MT



Construction of t-table



Sentence-aligned
corpus



Word alignments



cat chat 0.9
the cat le chat 0.8
dog chien 0.8
house maison 0.6
my house ma maison 0.9
language langue 0.9
...

Phrase table
(translation model)

Word Alignment Models

Lexical Translation

- How do we translate a word? Look it up in the dictionary

Haus – house, building, home, household, shell

- Multiple translations

- Some more frequent than others
 - Different word senses, different registers, different functions
 - *House, home* are common
- *Shell* is specialized (the Haus of a snail is a shell)

How Common is Each?

- Look at a parallel corpus (German text along with English translation)

Translation of Haus	Count
house	8000
building	1600
home	200
household	150
shell	50

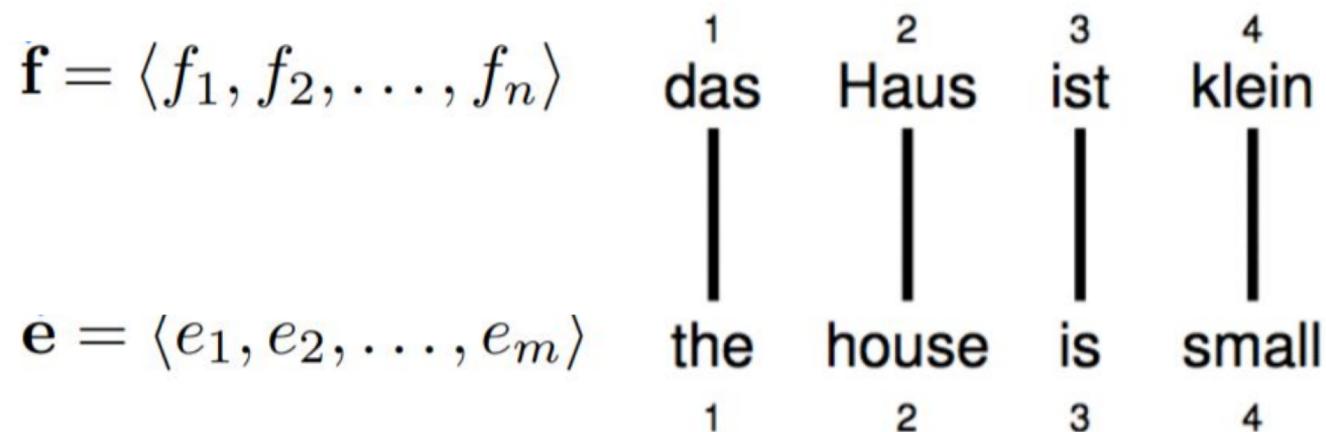
Estimate Translation Probabilities

- Maximum likelihood estimation

$$\hat{p}_{\text{MLE}}(e \mid \text{Haus}) = \begin{cases} 0.8 & \text{if } e = \text{house}, \\ 0.16 & \text{if } e = \text{building}, \\ 0.02 & \text{if } e = \text{home}, \\ 0.015 & \text{if } e = \text{household}, \\ 0.005 & \text{if } e = \text{shell}. \end{cases}$$

Word Alignment

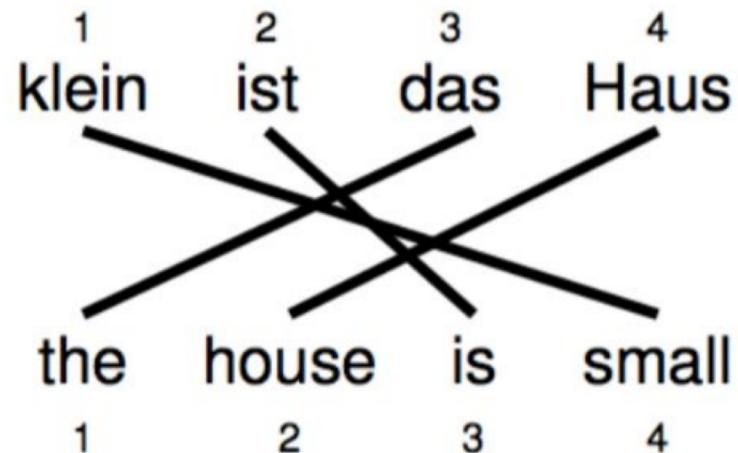
- Alignment can be visualized by drawing links between two sentences, and they are represented as vectors of positions



$$\mathbf{a} = (1, 2, 3, 4)^\top$$

Reordering

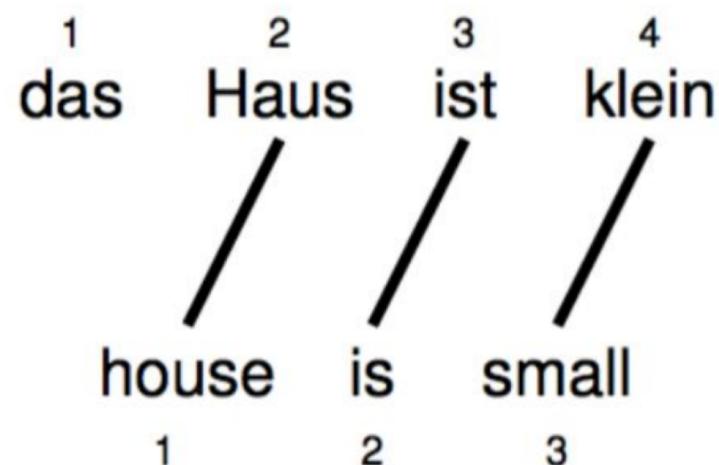
- Words may be reordered during translation



$$\mathbf{a} = (3, 4, 2, 1)^\top$$

Word Dropping

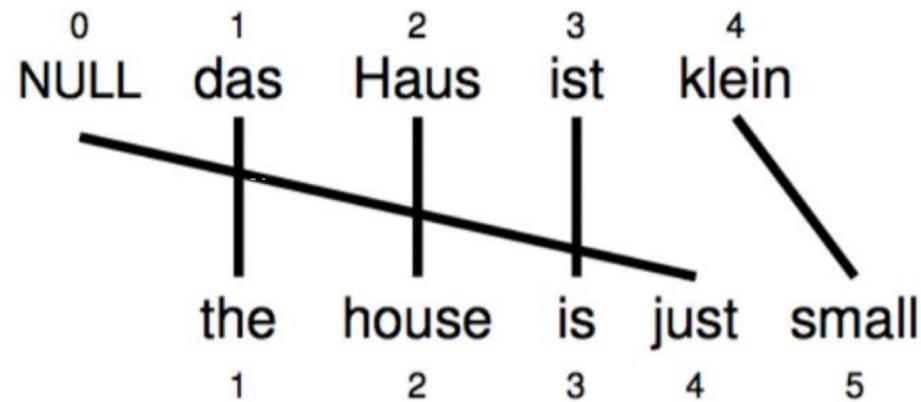
- A source word may not be translated at all



$$\mathbf{a} = (2, 3, 4)^\top$$

Word Insertion

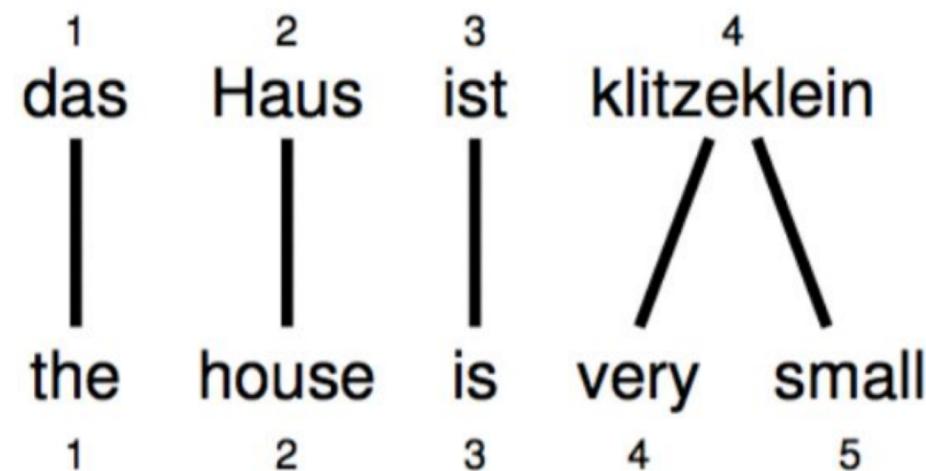
- Words may be inserted during translation
 - English **just** does not have an equivalent
 - But it must be explained – we typically assume every source sentence contains a **NULL** token



$$\mathbf{a} = (1, 2, 3, 0, 4)^{\top}$$

One-to-many Translation

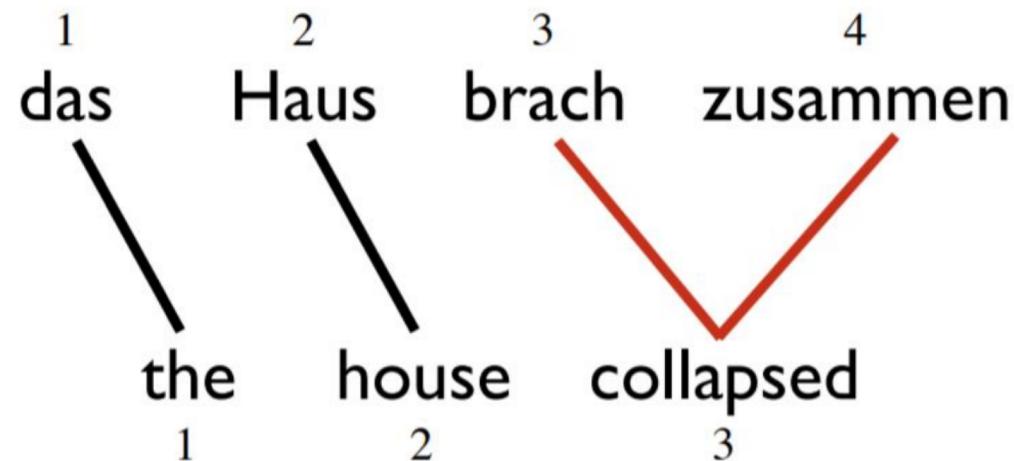
- A source word may translate into **more than one** target word



$$\mathbf{a} = (1, 2, 3, 4, 4)^\top$$

Many-to-one Translation

- More than one source word may **not** translate as a unit in lexical translation



$$\mathbf{a} = ???$$

$$\mathbf{a} = (1, 2, (3, 4)^\top)^\top ?$$

Computing Word Alignments

- Word alignments are the basis for most translation algorithms
- Given two sentences F and E, find a good alignment
- But a word-alignment algorithm can also be part of a mini-translation model itself
- One the most basic alignment models is also a simplistic translation model

IBM Model 1

- Generative model: break up translation process into smaller steps
- Simplest possible lexical translation model
- Additional assumptions
 - All alignment decisions are independent
 - The alignment distribution for each a_i is uniform over all source words and NULL

Lexical Translation

- Goal: a model $p(\mathbf{e}|\mathbf{f}, m)$
 - Where **e** and **f** are complete English and Foreign sentences
 - $\mathbf{e} = \langle e_1, e_2, \dots, e_m \rangle$
 - $\mathbf{f} = \langle f_1, f_2, \dots, f_n \rangle$

Lexical Translation

- Goal: a model $p(\mathbf{e}|\mathbf{f}, \mathbf{m})$
 - Where **e** and **f** are complete English and Foreign sentences
- Lexical translation makes the following assumptions
 - Each word e_i in e is generated from exactly one word in f
 - Thus, we have an alignment a_i that indicates which word e_i “came from”, specifically it came from f_{a_i}
 - Given the alignments **a**, translation decisions are conditionally independent of each other and depend only on the aligned source word f_{a_i}

Lexical Translation

- Putting our assumptions together, we have:

$$p(\mathbf{e}|\mathbf{f}, \mathbf{m}) = \sum_{\mathbf{a} \in [0,n]^m} p(\mathbf{a}|\mathbf{f}, \mathbf{m}) \times \prod_{i=1}^m p(e_i|f_{a_i})$$

Alignment × Translation | Alignment

IBM Model 1: $P(E|F)$

- Translation probability
 - For a foreign sentence $f = (f_1, \dots, f_{l_f})$ of length l_f
 - To an English sentence $e = (e_1, \dots, e_{l_e})$ of length l_e
 - With an alignment of each English word e_j to a foreign word f_i according to the alignment function $a: j \rightarrow i$

$$p(e, a | f) = \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)})$$

- Parameter ϵ is a normalization constant

Computing $P(E|F)$ in IBM Model 1

$$p(e, a | f) = \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)})$$

- A normalization factor, since there are $(l_f + 1)^{l_e}$ possible alignments
- Parameter ϵ is a normalization constant
- The probability of an alignment given the foreign sentence

Computing $P(E|F)$ in IBM Model 1

$$p(e, a | f) = \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)})$$

$$p(e|f) = \sum_a p(e, a | f) = \sum_a p(a|f) \times \prod_{j=1}^{l_e} p(e_j | f_{a_j})$$

Example

das	
<i>e</i>	$t(e f)$
the	0.7
that	0.15
which	0.075
who	0.05
this	0.025

Haus	
<i>e</i>	$t(e f)$
house	0.8
building	0.16
home	0.02
household	0.015
shell	0.005

ist	
<i>e</i>	$t(e f)$
is	0.8
's	0.16
exists	0.02
has	0.015
are	0.005

klein	
<i>e</i>	$t(e f)$
small	0.4
little	0.4
short	0.1
minor	0.06
petty	0.04

$$\begin{aligned}
 p(e, a|f) &= \frac{\epsilon}{4^3} \times t(\text{the}| \text{das}) \times t(\text{house}| \text{Haus}) \times t(\text{is}| \text{ist}) \times t(\text{small}| \text{klein}) \\
 &= \frac{\epsilon}{4^3} \times 0.7 \times 0.8 \times 0.8 \times 0.4 \\
 &= 0.0028\epsilon
 \end{aligned}$$

Estimate Translation Probabilities

- Maximum likelihood estimation

$$\hat{p}_{\text{MLE}}(e \mid \text{Haus}) = \begin{cases} 0.8 & \text{if } e = \text{house}, \\ 0.16 & \text{if } e = \text{building}, \\ 0.02 & \text{if } e = \text{home}, \\ 0.015 & \text{if } e = \text{household}, \\ 0.005 & \text{if } e = \text{shell}. \end{cases}$$

Estimate Alignments Given t-table

- If we have translation probabilities

<i>e</i>	$t(e f)$
the	0.7
that	0.15
which	0.075
who	0.05
this	0.025

<i>e</i>	$t(e f)$
house	0.8
building	0.16
home	0.02
household	0.015
shell	0.005

<i>e</i>	$t(e f)$
is	0.8
's	0.16
exists	0.02
has	0.015
are	0.005

<i>e</i>	$t(e f)$
small	0.4
little	0.4
short	0.1
minor	0.06
petty	0.04

- The goal is to find the most probable alignment given a parameterized model

Estimating the Alignment

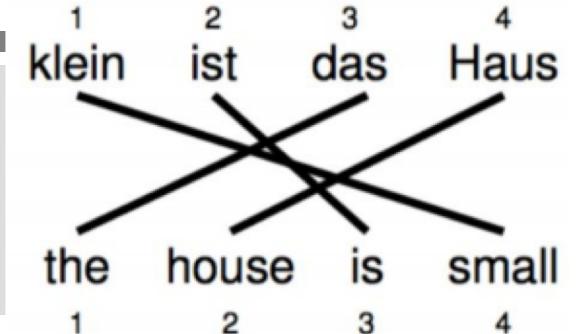
$$p(e, a | f) = \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)})$$

$$\begin{aligned} a^* &= \arg \max_a p(e, a | f) \\ &= \arg \max_a \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)}) \\ &= \arg \max_a \prod_{j=1}^{l_e} t(e_j | f_{a(j)}) \end{aligned}$$

Since translation choice for each position is independent, the product is maximized by maximizing each term:

$$a_i^* = \arg \max_{a_i=0}^n t(e_i | f_{a_i})$$

Learning Lexical Translation Models



- We'd like to estimate the lexical translation probabilities $t(e|f)$ from a parallel corpus but we do not have the alignments
- **Chick and egg problem**
 - If we had the **alignments**, we could estimate the parameters of our generative model (MLE)
 - If we had the **parameters**, we could estimate the alignments

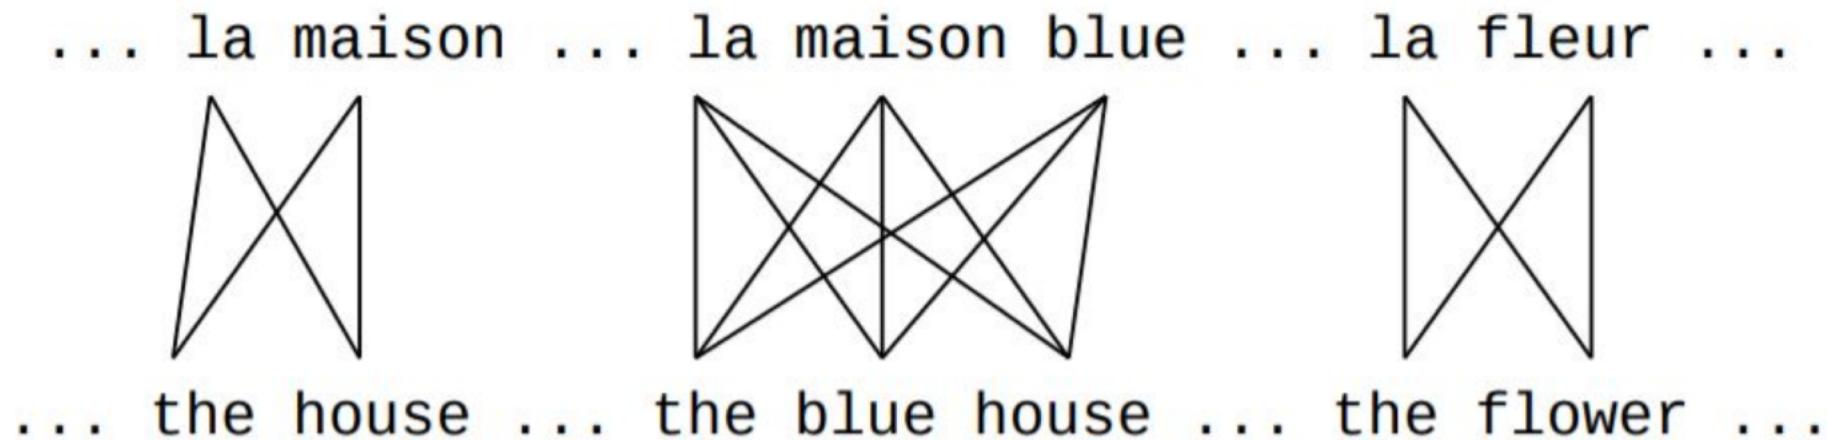
klein	
<i>e</i>	$t(e f)$
small	0.4
little	0.4
short	0.1
minor	0.06
petty	0.04

EM Algorithm

- Incomplete data
 - If we had **complete data**, we could estimate the model
 - If we had the **model**, we could fill in the gaps in the data
- **Expectation Maximization (EM) in a nutshell**
 1. Initialize model parameters (e.g., uniform, random)
 2. Assign probabilities to the missing data
 3. Estimate model parameters from completed data
 4. Iterate steps 2-3 until convergence

EM Algorithm

Kevin Knight's example



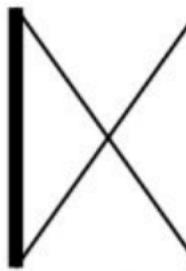
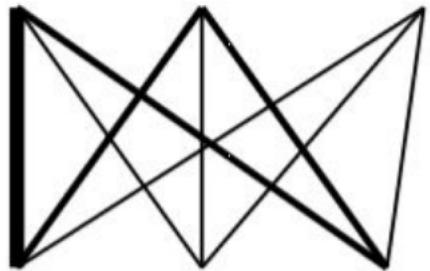
- Initial step: all word alignments equally likely
- Model learns that: e.g., *la* is often aligned with *the*

EM Algorithm

... la maison ... la maison blue ... la fleur ...

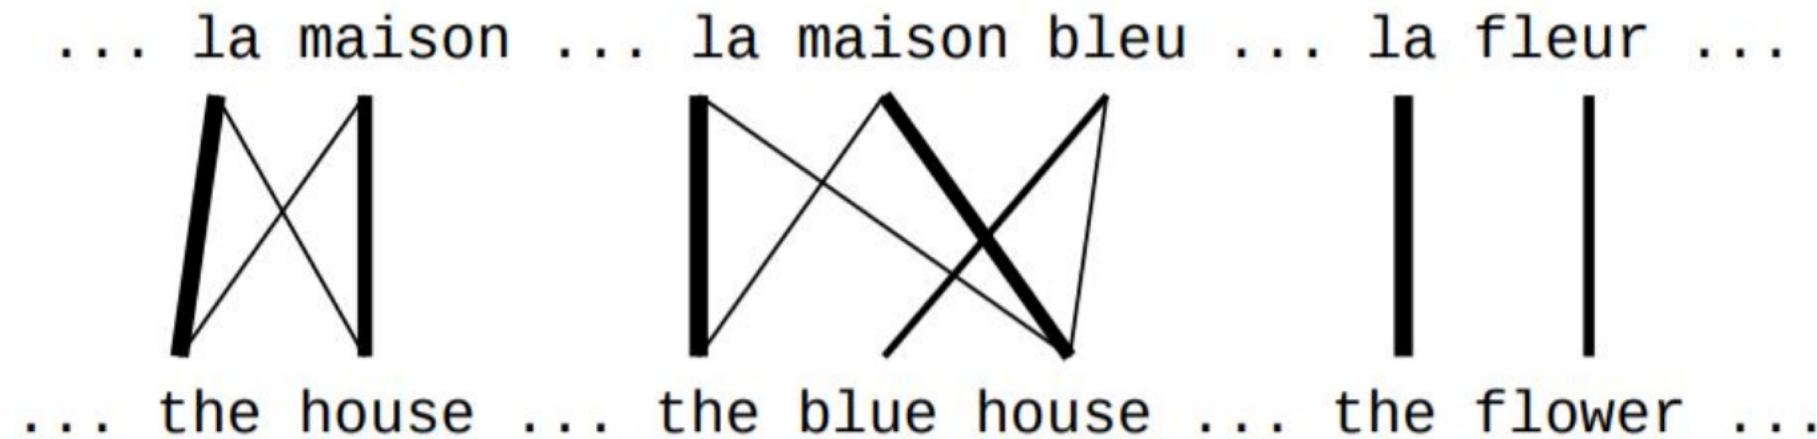


... the house ... the blue house ... the flower ...



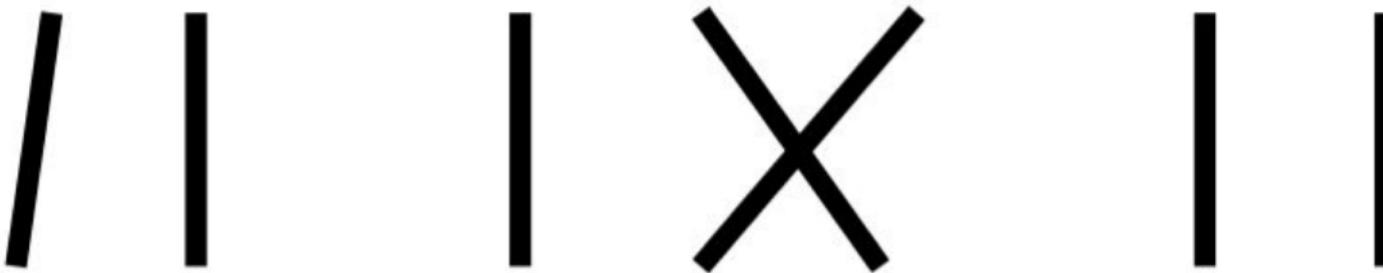
- After one iteration
- Alignments, e.g., between *la* and *the* are more likely

EM Algorithm



- After another iteration
- It becomes apparent that alignments, e.g., between *fleur* and *flower* are more likely

EM Algorithm

... la maison ... la maison bleu ... la fleur ...

... the house ... the blue house ... the flower ...

- Convergence
- Inherent hidden structure revealed by EM !

EM Algorithm

... la maison ... la maison bleu ... la fleur ...



... the house ... the blue house ... the flower ...



$$\begin{aligned} p(\text{la}|\text{the}) &= 0.453 \\ p(\text{le}|\text{the}) &= 0.334 \\ p(\text{maison}|\text{house}) &= 0.876 \\ p(\text{bleu}|\text{blue}) &= 0.563 \\ \dots \end{aligned}$$

- Parameter estimation from the aligned corpus

The EM Algorithm for Word Alignment

- Initialize the mode, typically with uniform distributions
- Repeat
 - **E Step:** use the current model to compute the probability of all possible alignments of the training data
 - **M Step:** use these alignment probability estimates to re-estimate values for all of the parameters
- Until convergence (i.e., parameters no longer change)

IBM Model 1 and EM

t-table Probabilities

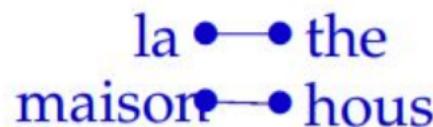
$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

IBM Model 1 and EM

t-table Probabilities

$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

Alignments

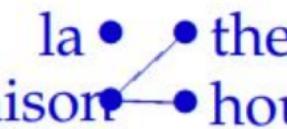
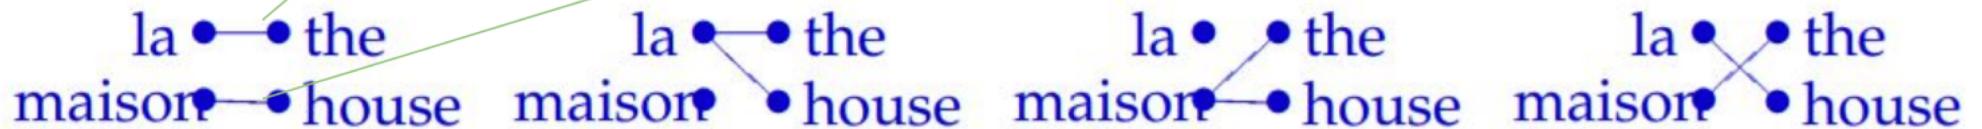


IBM Model 1 and EM

t-table Probabilities

$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

Alignments



$$p(\mathbf{e}, a|\mathbf{f}) = 0.56$$

$$p(\mathbf{e}, a|\mathbf{f}) = 0.035$$

$$p(\mathbf{e}, a|\mathbf{f}) = 0.08$$

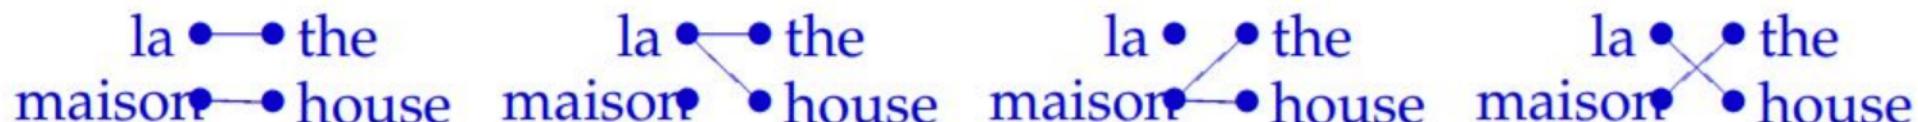
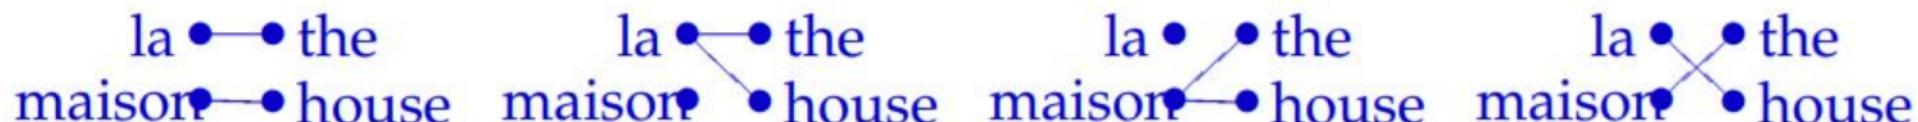
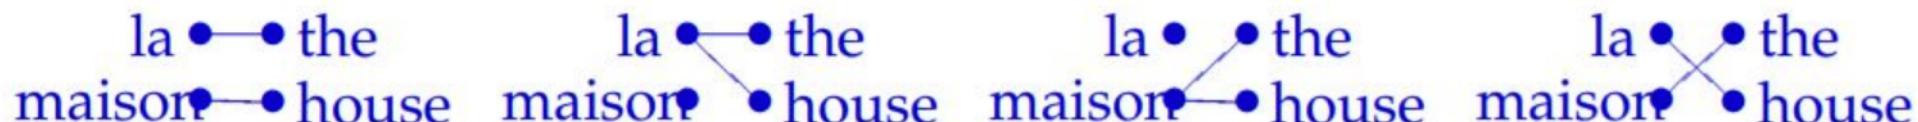
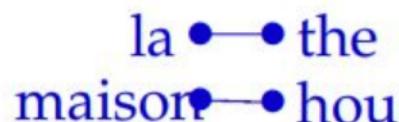
$$p(\mathbf{e}, a|\mathbf{f}) = 0.005$$

IBM Model 1 and EM

t-table Probabilities

$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

Alignments



$$p(\mathbf{e}, a | \mathbf{f}) = 0.56$$

$$p(\mathbf{e}, a | \mathbf{f}) = 0.035$$

$$p(\mathbf{e}, a | \mathbf{f}) = 0.08$$

$$p(\mathbf{e}, a | \mathbf{f}) = 0.005$$

$$\text{Chain rule: } p(a | \mathbf{e}, \mathbf{f}) = \frac{p(\mathbf{e}, a | \mathbf{f})}{p(\mathbf{e} | \mathbf{f})}$$

IBM Model 1 and EM: Expectation Step

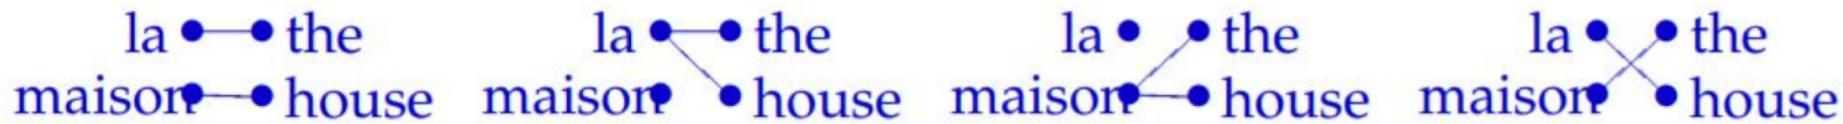
$$\begin{aligned} p(a | \mathbf{e}, \mathbf{f}) &= \frac{p(\mathbf{e}, a | \mathbf{f})}{p(\mathbf{e} | \mathbf{f})} \\ &= \frac{\frac{\epsilon}{(l_f+1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a_j})}{\frac{\epsilon}{(l_f+1)^{l_e}} \prod_{j=1}^{l_e} \sum_{i=0}^{l_f} t(e_j | f_{a_j})} \\ &= \prod_{j=1}^{l_e} \frac{t(e_j | f_{a_j})}{\sum_{i=0}^{l_f} t(e_j | f_{a_j})} \end{aligned}$$

IBM Model 1 and EM

t-table Probabilities

$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

Alignments



$$p(\mathbf{e}, a | \mathbf{f}) = 0.56 \quad p(\mathbf{e}, a | \mathbf{f}) = 0.035 \quad p(\mathbf{e}, a | \mathbf{f}) = 0.08 \quad p(\mathbf{e}, a | \mathbf{f}) = 0.005$$

E-step

$$p(a|\mathbf{e}, \mathbf{f}) = 0.824 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.052 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.118 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.007$$

IBM Model 1 and EM: Maximization Step

- Now we have to collect counts
- Evidence from a sentence pair (e, f) that word e is a translation of word f :

$$c(e|f) = \sum_a p(a|e, f) \sum_{j=1}^{l_e} \delta(e, e_j) \delta(f, f_{a_j})$$

IBM Model 1 and EM

t-table Probabilities

$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

Alignments

$\text{la} \bullet \text{---} \bullet \text{the}$ $\text{maison} \bullet \text{---} \bullet \text{house}$	$\text{la} \bullet \text{---} \bullet \text{the}$ $\text{maison} \bullet \text{---} \bullet \text{house}$	$\text{la} \bullet \text{---} \bullet \text{the}$ $\text{maison} \bullet \text{---} \bullet \text{house}$	$\text{la} \bullet \text{---} \bullet \text{the}$ $\text{maison} \bullet \text{---} \bullet \text{house}$
$p(\mathbf{e}, \mathbf{a} \mathbf{f}) = 0.56$	$p(\mathbf{e}, \mathbf{a} \mathbf{f}) = 0.035$	$p(\mathbf{e}, \mathbf{a} \mathbf{f}) = 0.08$	$p(\mathbf{e}, \mathbf{a} \mathbf{f}) = 0.005$

E-step

$$p(a|\mathbf{e}, \mathbf{f}) = 0.824 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.052 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.118 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.007$$

M-step Counts

$$\begin{array}{ll} c(\text{the}|\text{la}) = 0.824 + 0.052 & c(\text{house}|\text{la}) = 0.052 + 0.007 \\ c(\text{the}|\text{maison}) = 0.118 + 0.007 & c(\text{house}|\text{maison}) = 0.824 + 0.118 \end{array}$$

IBM Model 1 and EM

t-table **Probabilities**

$$\begin{array}{ll} p(\text{the}|\text{la}) = 0.7 & p(\text{house}|\text{la}) = 0.05 \\ p(\text{the}|\text{maison}) = 0.1 & p(\text{house}|\text{maison}) = 0.8 \end{array}$$

E-step **Alignments**

$$p(a|\mathbf{e}, \mathbf{f}) = 0.824 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.052 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.118 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.007$$

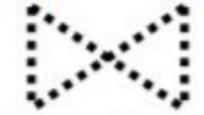
M-step **Counts**

$$\begin{array}{ll} c(\text{the}|\text{la}) = 0.824 + 0.052 & c(\text{house}|\text{la}) = 0.052 + 0.007 \\ c(\text{the}|\text{maison}) = 0.118 + 0.007 & c(\text{house}|\text{maison}) = 0.824 + 0.118 \end{array}$$

Update t-table:

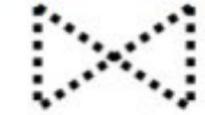
$$p(\text{the}|\text{la}) = c(\text{the}|\text{la})/c(\text{la})$$

Convergence

das Haus

the house

das Buch

the book

ein Buch

a book

e	f	initial	1st it.	2nd it.	3rd it.	...	final
the	das	0.25	0.5	0.6364	0.7479	...	1
book	das	0.25	0.25	0.1818	0.1208	...	0
house	das	0.25	0.25	0.1818	0.1313	...	0
the	buch	0.25	0.25	0.1818	0.1208	...	0
book	buch	0.25	0.5	0.6364	0.7479	...	1
a	buch	0.25	0.25	0.1818	0.1313	...	0
book	ein	0.25	0.5	0.4286	0.3466	...	0
a	ein	0.25	0.5	0.5714	0.6534	...	1
the	haus	0.25	0.5	0.4286	0.3466	...	0
house	haus	0.25	0.5	0.5714	0.6534	...	1

Word Alignment:

- Given a sentence pair, which words correspond to each other?

Michael	Michael	geht	davon	aus	,	dass	er	im	haus	bleibt
assumes										
that										
he										
will										
stay										
in										
the										
house										

Higher IBM Models

IBM Model 1	Lexical translation
IBM Model 2	Adds absolute reordering model
IBM Model 3	Adds fertility model
IBM Model 4	Relative reordering model
IBM Model 5	Fixes deficiency

- Only IBM Model 1 has global maximum
 - Training of a higher IBM model builds upon previous model
- Computationally biggest change in Model 3

Word Alignment and IBM Models

- IBM models create a many-to-one mapping
 - Words are aligned using an alignment function
 - A function may return the same value for different input (one-to-many mapping)
 - A function cannot return multiple values for one input (no many-to-one mapping)
- Real world alignments have many-to-many mappings

Evaluation Metrics

- Manual evaluation is most accurate, but expensive
- Automated evaluation metrics:
 - Compare system hypothesis with reference translations
 - BiLingual Evaluation Understudy (BLEU) (Papineni et al., 2002):
 - Modified n-gram precision

$$p_n = \frac{\text{number of } n\text{-grams appearing in both reference and hypothesis translations}}{\text{number of } n\text{-grams appearing in the hypothesis translation}}$$

BLEU

$$\text{BLEU} = \exp \frac{1}{N} \sum_{n=1}^N \log p_n$$

- Two modifications:
 - To avoid $\log 0$, all precisions are smoothed
 - Each n-gram in reference can be used at most once
 - Ex. **Hypothesis:** to to to to to vs **Reference:** to be or not to be should not get a unigram precision of 1
- Precision-based metrics favor short translations
 - Solution: Multiply score with a brevity penalty (BP) for translations shorter than reference, $e^{1-r/h}$

BLEU Scores

Translation		p_1	p_2	p_3	p_4	BP	BLEU
Reference	<i>Vinay likes programming in Python</i>						
Sys1	<i>To Vinay it like to program Python</i>	$\frac{2}{7}$	0	0	0	1	.21
Sys2	<i>Vinay likes Python</i>	$\frac{3}{3}$	$\frac{1}{2}$	0	0	.51	.33
Sys3	<i>Vinay likes programming in his pajamas</i>	$\frac{4}{6}$	$\frac{3}{5}$	$\frac{2}{4}$	$\frac{1}{3}$	1	.76

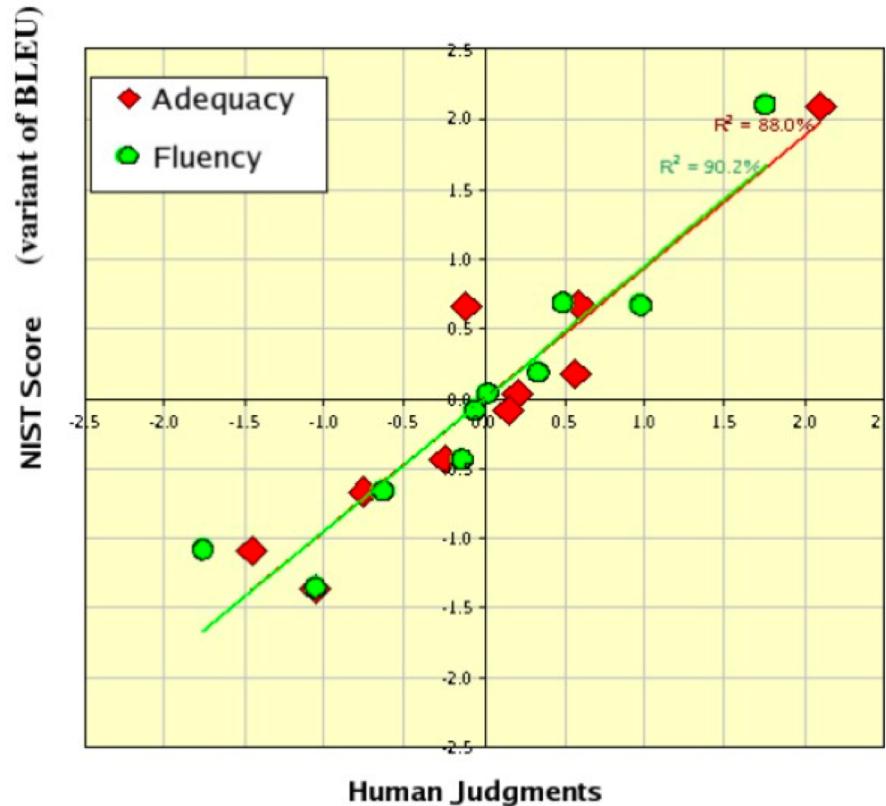
Sample BLEU scores for various system outputs

- Alternatives have been proposed:
 - METEOR: weighted F-measure
 - Translation Error Rate (TER): Edit distance between hypothesis and reference

Other Issues?

BLEU

- Correlates somewhat well with human judgments



(G. Doddington, NIST)

Problems with Lexical Translation

- Complexity – exponential in sentence length
- Weak reordering – the output is not fluent
- Many local decisions – error propagation